

# Isolating Performance Issues: A Filesystem Congestion

---

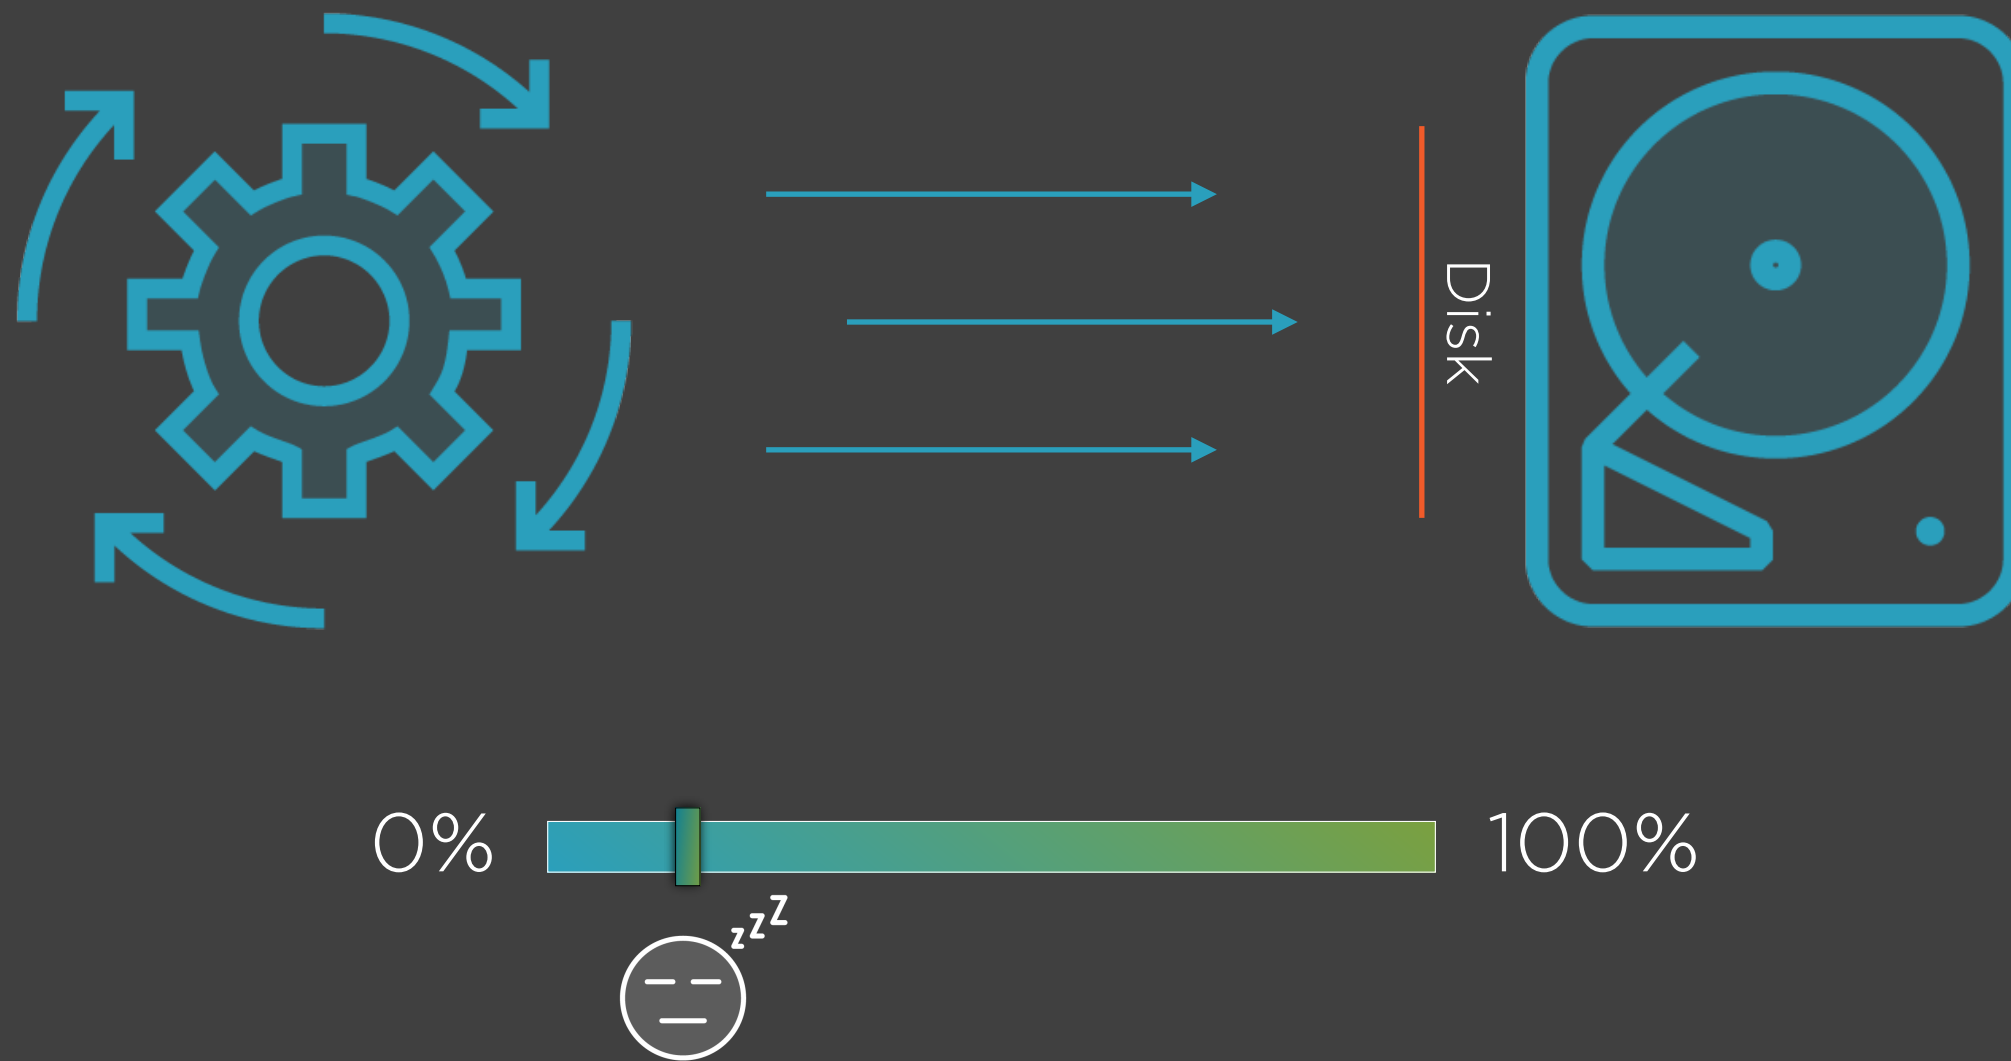


Uriah Levy

SOFTWARE ENGINEER

@iamuriah1 [www.medium.com/@iamuriah1](https://www.medium.com/@iamuriah1)

# An I/O-bound Application




# Filesystems Are Fast

Logging

Copying files from one location to another

Downloading files from the internet



# Overview

**Understanding Filesystem read/write mechanics in Java**

**Identifying a filesystem bottleneck while reading/writing files using thread dumps**

**Identifying “Thread Stagnation”**

**Examining the application’s resource-footprint**

# Filesystem Read/Write Mechanics in Java

---

# Input/Output Streams

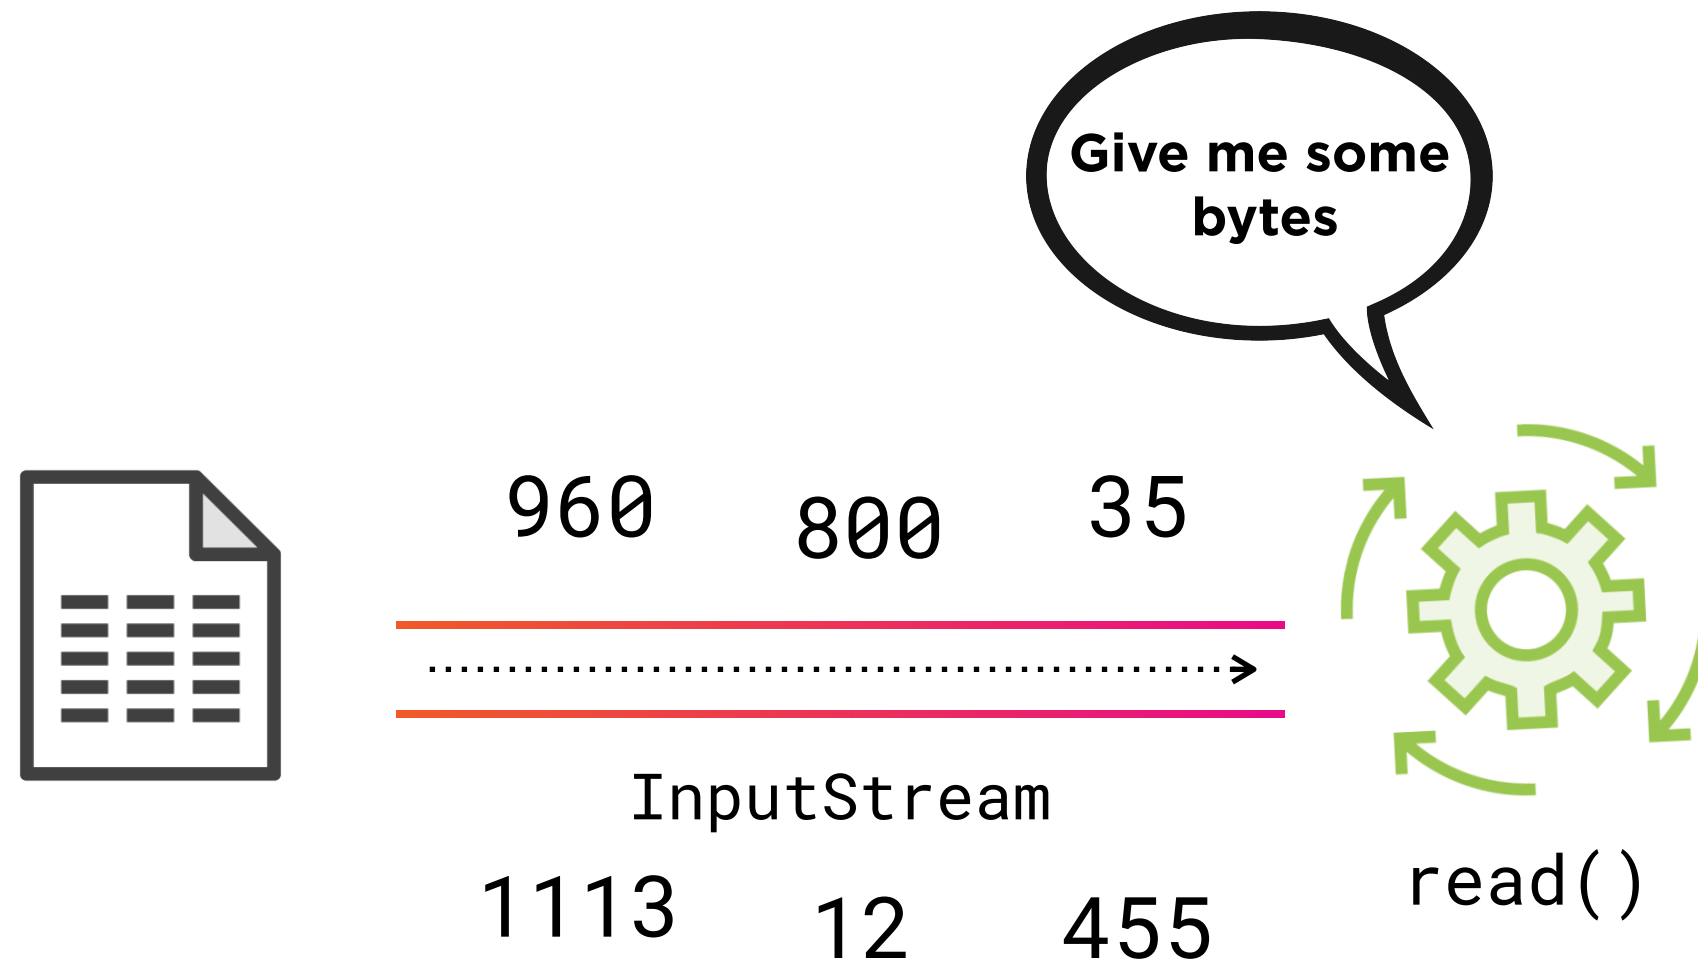
**InputStream:**

**Reads data**

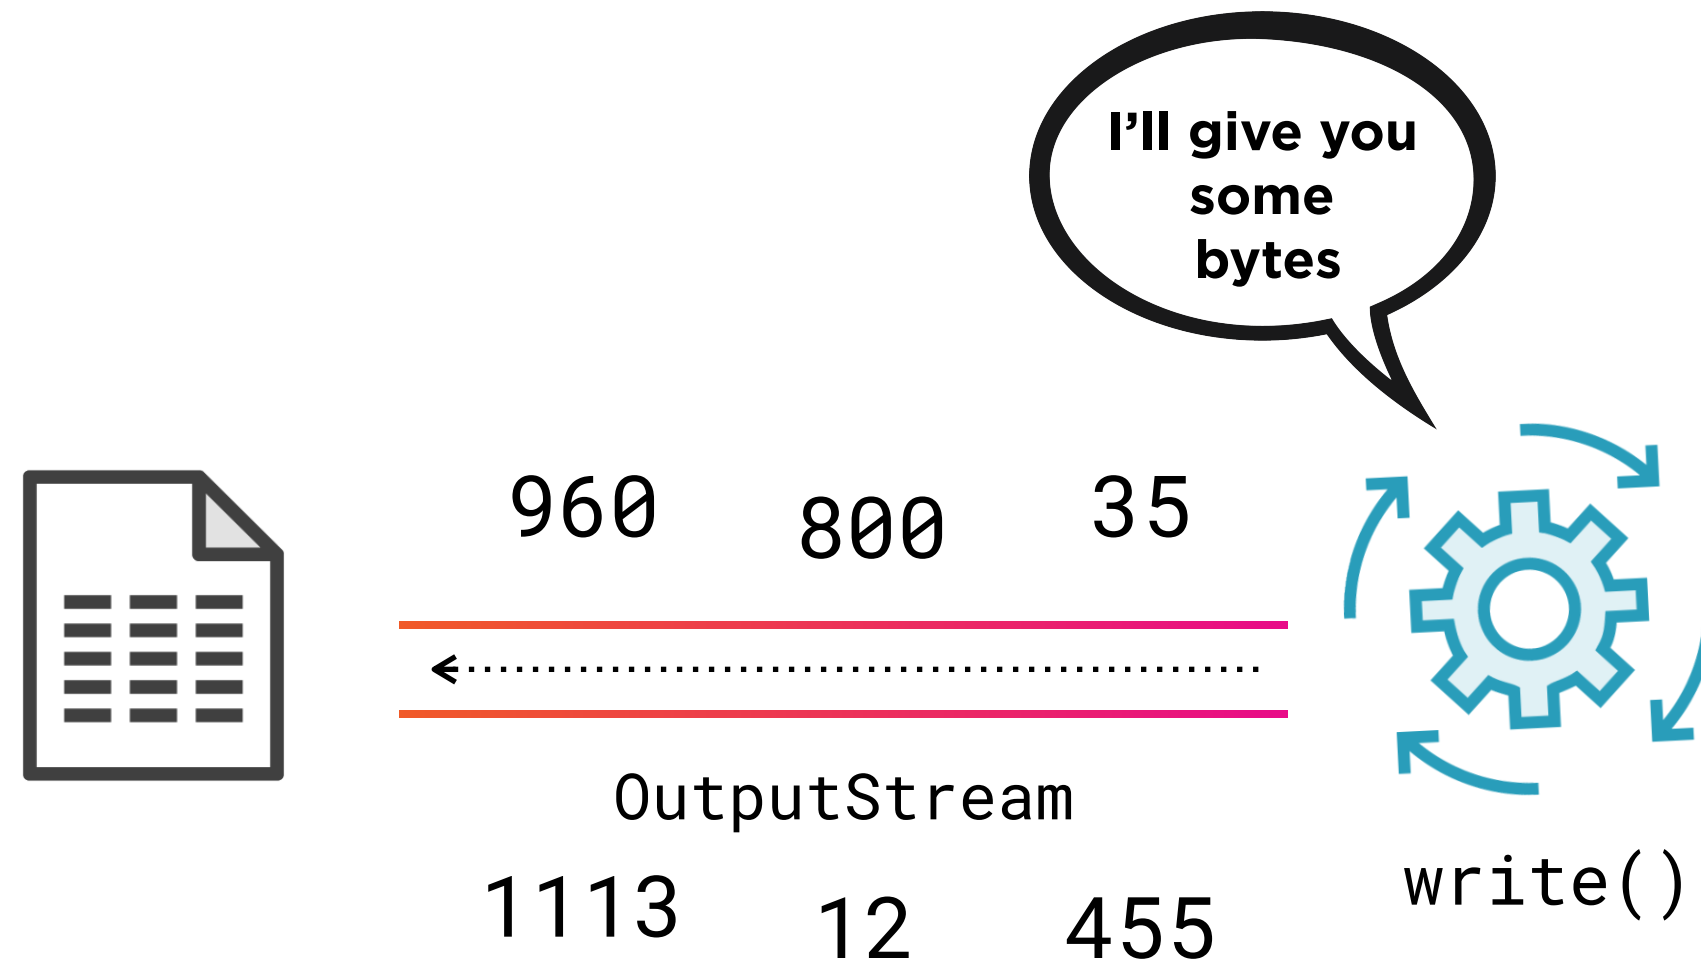
**OutputStream:**

**Writes data**

# Reading from a Filesystem

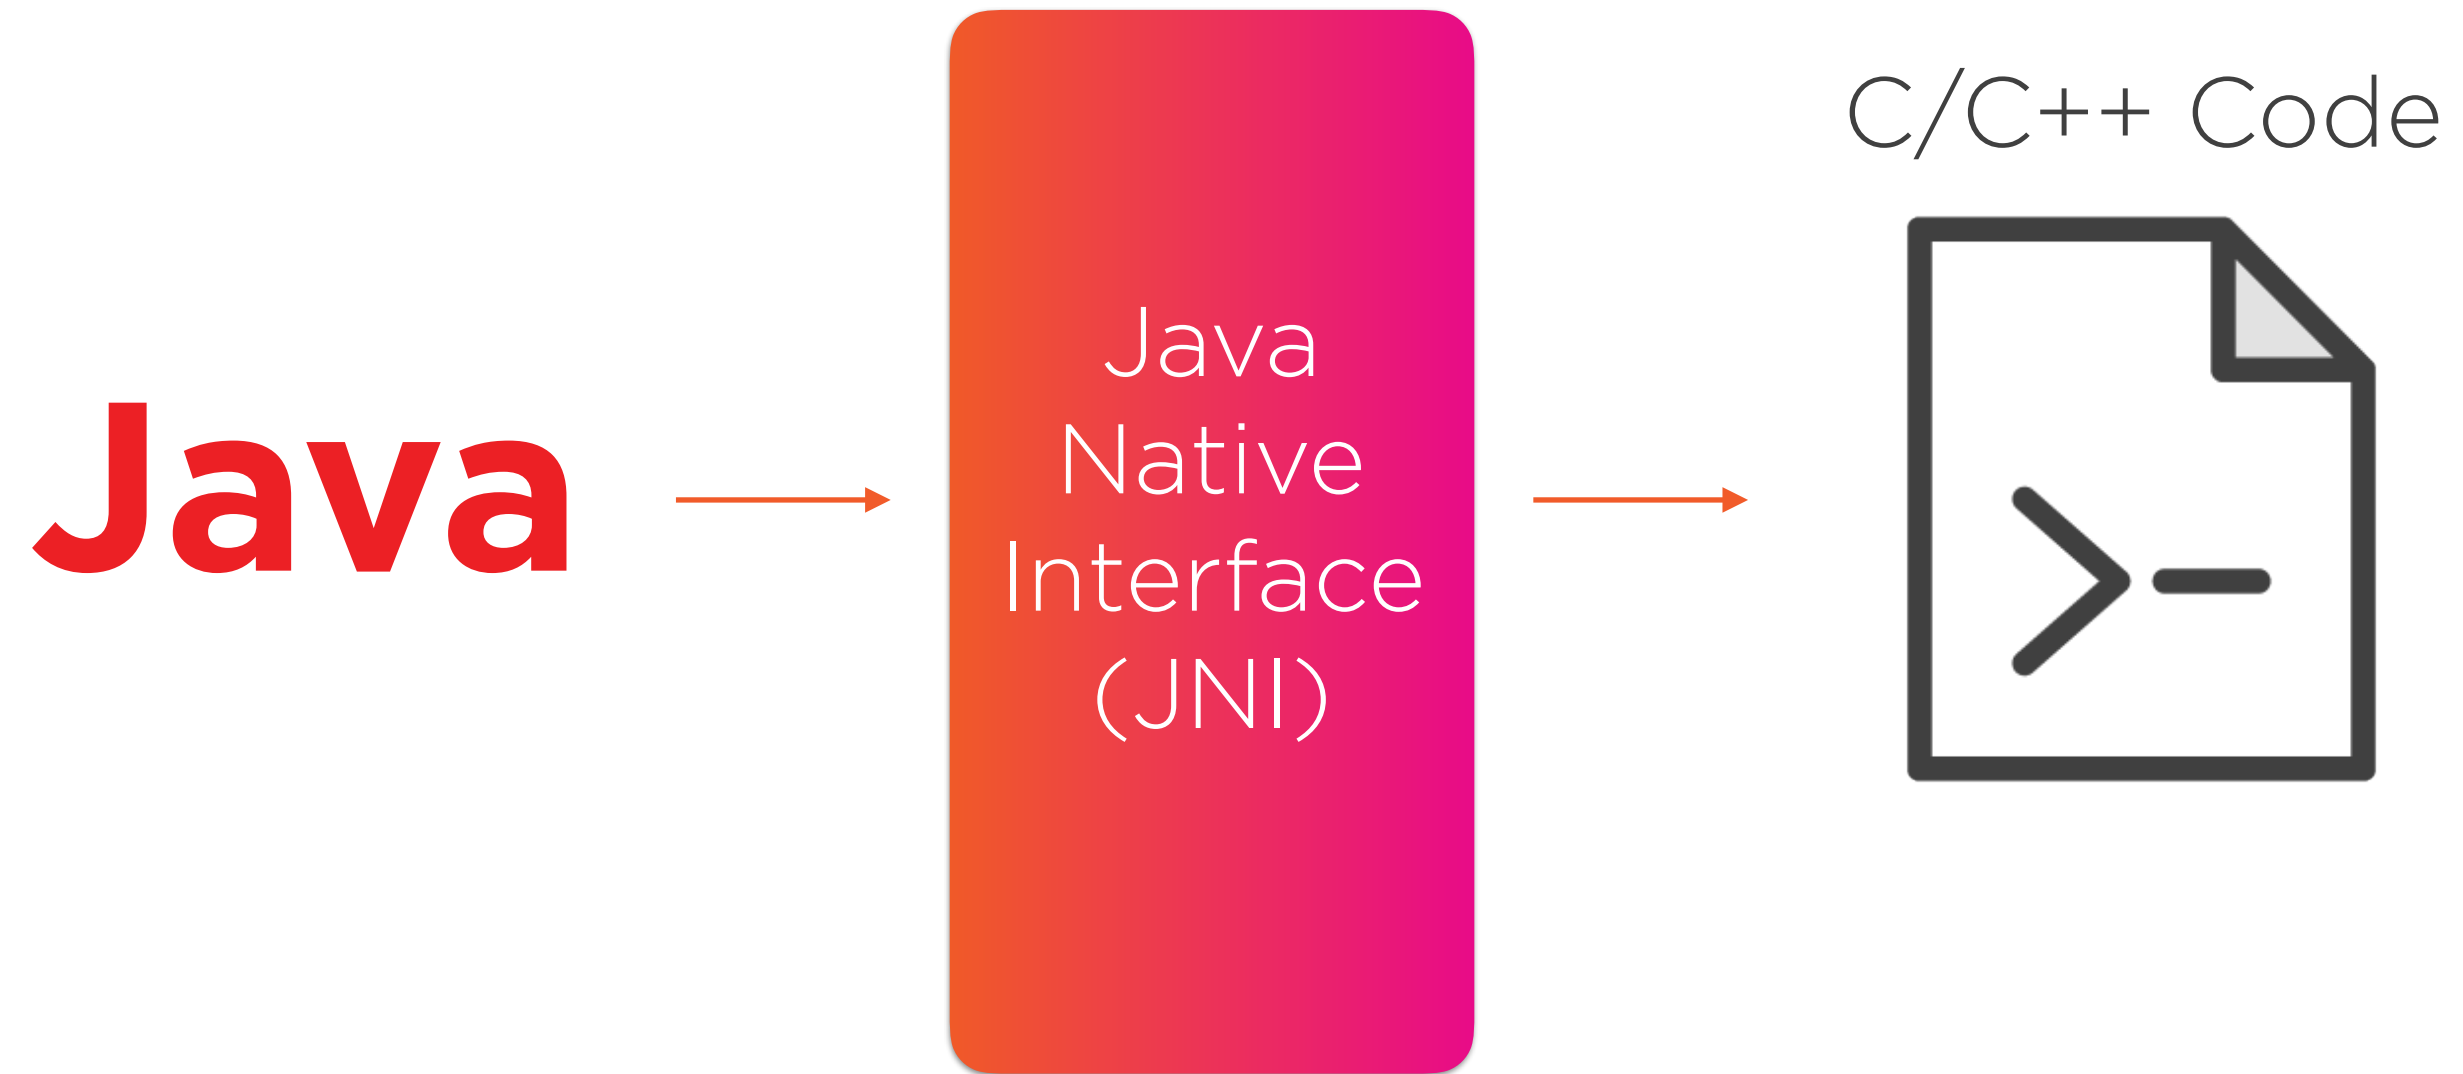


# Writing to a Filesystem





# Low-level Action Delegation



Remember Native Methods?

```
"pool-1-thread-1" #9 prio=5 os_prio=0 tid=0x00007f00001e9800 nid=0x6b15 runnable [0x00007effe9147000]  
  java.lang.Thread.State: RUNNABLE  
    ↪ at net.smacke.jaydio.DirectIoLib.pwrite(Native Method)  
      at net.smacke.jaydio.DirectIoLib.pwrite(DirectIoLib.java:223)  
        at net.smacke.jaydio.channel.DirectIoByteChannel.write(DirectIoByteChannel.java:93)  
          at net.smacke.jaydio.channel.DirectIoByteChannel.write(DirectIoByteChannel.java:41)  
            at net.smacke.jaydio.align.ByteChannelAligner.flush(ByteChannelAligner.java:321)  
              at net.smacke.jaydio.align.ByteChannelAligner.truncate(ByteChannelAligner.java:292)  
                at net.smacke.jaydio.align.ByteChannelAligner.close(ByteChannelAligner.java:88)  
                  at net.smacke.jaydio.DirectRandomAccessFile.close(DirectRandomAccessFile.java:97)  
                    at org.uriah1.ajtd.filesystem.DirectIoWriterCallable.call(FilesystemOperator.java:136)
```

---


## Native Method in the Stack

# Docker / Linux cgroups



Docker

+



cgroups

```
docker run -it --device-read-bps /dev/sda:150kb ajtd-  
ubuntu:latest /bin/bash
```

---

Spawns a container with **reads** limited to  
150KB per-second

```
docker run -it --device-write-bps /dev/sda:150kb ajtd-ubuntu:latest /bin/bash
```

---

Spawns a container with **writes** limited to 150KB per-second

# The “blkio” Controller

```
> cat /sys/fs/cgroup/blkio/blkio.throttle.read_bps_device
```

```
8:0 153600
```

```
> cat /sys/fs/cgroup/blkio/blkio.throttle.write_bps_device
```

```
8:0 153600
```

## > Demo

---

Follow the README.txt from the course assets to be able to follow along on your environment



<https://github.com/smacke/jaydio>

# Demo

**Identifying “Thread Stagnation”**

**Debunking Thread Stagnation using  
‘strace’**

**Examining the application’s resource-  
footprint**

<https://linux.die.net/man/1/strace>

\* Follow the README file from the Module 4 assets to deploy our strace-equipped docker container locally

# Debugging Chain Pseudo-Snippet

First, invoke fs-operator:

- \*while fs-operator is running:

- Take some thread dumps

- Attach to the JVM process with 'strace'

Invoke 'fs-operator.jar'

```
java -jar fs-operator.jar read 1 | ./nid-translator.pl >  
thread-dumps-w-decimal-nids.out &
```

# Translate 'nid' from Hex to Decimal

```
1  #!/usr/bin/perl -w
2
3  while (<>) {
4      if (/nid=(0x[[:xdigit:]]+)/) {
5          $lwp = hex($1);
6          s/nid=/lwp=$lwp nid=/;
7      }
8      print;
9  }
```

# Capture 10 Thread Dumps

```
1  #!/bin/bash
2
3  a=1
4  numberOfKills=10
5  pid=$(ps aux | grep [j]ava | awk '{print $2}')
6
7  echo "Sending kill -3 at PID:" $pid
8
9  while [ $a -le $numberOfKills ]; do
10      kill -3 $pid
11      sleep 1
12      a=`expr $a + 1`
13  done
```

# Attach with 'strace'

Get the PID programatically

Specify the output filename(s) prefix

---

```
ps aux | grep [j]ava | awk '{print $2}' | xargs strace -o  
fs-operator -tt -T -ff -p
```

Indicate system time in microseconds

Trace child processes

Indicate total call time



# The Entire Debugging Chain

```
1  #! /bin/bash
2
3  java -jar fs-operator.jar read 1 | ./nid-translator.pl
4  > thread-dumps-w-decimal-nids.out &
5
6  sh dump-threads.sh &
7
8  ps aux | grep [j]ava | awk '{print $2}' | xargs strace -o
9  fs-operator -tt -T -ff -p
```